

Course Overview

CS-392-A Systems Programming

Instructors:

Georgios Portokalidis (*This section*)

Srinivas Sridharan (Other section)

Updated textbook and content!

Based on CMU's Computer Systems course

Overview

- **Course theme**
- **Five realities**
- **Academic integrity**
- **Logistics**

Course Theme:

Abstraction Is Good But Don't Forget Reality

- **Most CS courses emphasize abstraction**

- Abstract data types
- Asymptotic analysis

- **These abstractions have limits**

- Especially in the presence of bugs
- Need to understand details of underlying implementations

Course Theme:

Useful outcomes from this course

- **Become more effective programmers**
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance


- **Prepare for later “systems” classes in CS**
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!



■ Int's:

- $40000 * 40000 \approx 1600000000$
- $50000 * 50000 \approx ?$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0)	⌘	3.14
fun(1)	⌘	3.14
fun(2)	⌘	3.1399998664856
fun(3)	⌘	2.00000061035156
fun(4)	⌘	3.14
fun(6)	⌘	Segmentation fault

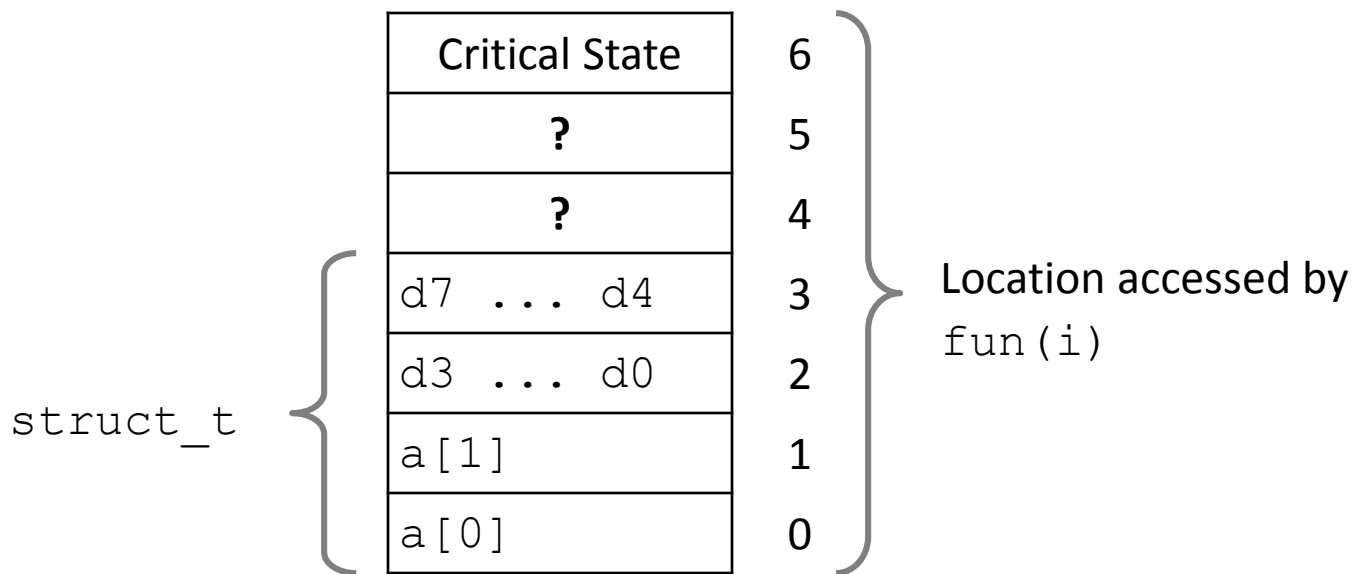
- Result is system specific

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

fun(0)	⌘	3.14
fun(1)	⌘	3.14
fun(2)	⌘	3.1399998664856
fun(3)	⌘	2.00000061035156
fun(4)	⌘	3.14
fun(6)	⌘	Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g., Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```


4.3ms

2.0 GHz Intel Core i7 Haswell

81.8ms

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance

- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

■ Most Systems Courses are Builder-Centric

- Computer Architecture
 - Design pipelined processor in Verilog
- Operating Systems
 - Implement sample portions of operating system
- Compilers
 - Write compiler for simple language
- Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**

Staff for This Section

■ Instructor: Georgios Portokalidis (that's me)

- Email: Georgios.Portokalidis@stevens.edu

■ Course Assistants

- Nicholas Cyprus (ncyprus@stevens.edu)
- Benjamin Iofel (biofel@stevens.edu)

Cheating: Description

■ Please pay close attention

■ What is cheating?

- Sharing code: by copying, retyping, **looking at**, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a lab, line by line
- Searching the Web for solutions
- Copying code from a previous course or online solution
 - You are only allowed to use code we supply, or from the CS:APP website

■ What is NOT cheating?

- Explaining how to use systems or tools
- Helping others with high-level design issues

■ Ignorance is not an excuse

Cheating: Consequences

■ Penalty for cheating:

- You will be reported to the Dean and the Honor board
- Penalties may include suspension and expulsion

■ Detection of cheating:

- We have sophisticated tools for detecting code plagiarism

■ Don't do it!

- Start early
 - You won't, but really start early!
- Ask us for help when you get stuck

Textbooks

■ Randal E. Bryant and David R. O'Hallaron,

- *Computer Systems: A Programmer's Perspective, Third Edition* (CS:APP3e), Pearson, 2016
- <http://csapp.cs.cmu.edu>
- This book really matters for the course!
 - How to solve labs
 - Practice problems typical of exam problems

■ Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators

■ Helpful but optional: W. Richard Stevens, Stephen A. Rago

- *Advanced Programming in the UNIX Environment*, 3rd ed., Addison-Wesley Professional, 2013

Course Components

■ Lectures (first 100 mins)

- Higher level concepts

■ Lectures (last 50 mins)

- Applied concepts, important tools and skills for labs, clarification of lectures, exam coverage
- In-class programming exercises
 - Make sure you have enough battery left

■ Quizzes (2 midterms)

- Test your understanding of concepts & mathematical principles

Course Components (cont)

■ Take-home assignments or labs (5)

- The heart of the course
- 2-3 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement
- Will be done individually

■ Final project

- Similar to labs
- Done at the end
- Small groups of maximum 2 students allowed
- Small demo and presentation of measurement results

Getting Help

■ Canvas

- Complete schedule of lectures, exams, and assignments
- Copies of lectures and assignments

■ Slack

- <https://cs392-2017s.slack.com/x-132211270262-132212553142/signup>
- Join with your Stevens email
- Discuss with your classmates
- Get feedback from CAs and instructors
- Do not post grading-related requests and grievances

■ Gradescope

- <https://gradescope.com/courses/6111>
- Submission of assignments
- Entry code 9X4BY9

Getting Help

■ Instructor Office hours:

- Tuesdays, 5:00pm-7:00pm, Lieb 214
- Try emailing first
- Appointments for meeting outside office hours may also possible
- Contact through email NOT canvas

■ CA Office hours

- Wednesdays, 4:00pm-6:00pm, Lieb lounge (3rd floor)
- Fridays, 3:00pm-5:00pm, Lieb lounge (3rd floor)
- Check Canvas for announcements
- Appointments for meeting outside office hours may also possible

Policies: Labs And Exams

■ Work groups

- You must work alone on all take-home assignments

■ Handins

- Labs due at 11:59pm on Tuesday
- Electronic handins using through gradescope
 - See assignment description for exactly what to download

■ Quizzes

- Quizzes will be during first part of lecture
- See syllabus for more info

■ Appealing grades

- By email me within 7 days from the posting of grades
- Appeal can lead to **higher** or **lower** grade

Development Platform

■ Use Xubuntu 16.04 x86-64 VM provided

- <https://www.dropbox.com/s/144ahwra1i6uaxk/cs392-xubuntu16.04.zip?dl=0>
- You can use your own platform but your submitted code must compile and execute correctly on this VM
- Submitted files must be as described in assignment
- You can run the VM using
 - Vmplayer – freely available from Vmware
 - VirtualBox – free software

Timeliness

■ Grace days

- **5 grace days** for the semester
- **Limit of 2 grace days** per assignment used **automatically**
- Covers scheduling crunch, out-of-town trips, illnesses, minor setbacks
- Save them until late in the term!

■ Lateness penalties

- Once grace day(s) used up, get penalized **15% per day**
- No handins later than **3 days after due date**

Timeliness (cont.)

■ Catastrophic events

- Major illness, death in family, ...
- Need to contact the office of undergraduate academics
 - They will contact me to make arrangements

■ Advice

- Once you start running late, it's really hard to catch up

Other Rules of the Lecture Hall

- **Laptops: permitted**

- But you need to make sure to have enough power for any in-class assignments at the end of the lecture

- **Electronic communications: *forbidden***

- No email, instant messaging, cell phone calls, etc

- **Presence in lectures: voluntary, recommended**

- **No recordings of ANY KIND**

Policies: Grading

- **Quizzes (30%): midterm I (15%), midterm II (15%)**
- **Take-home assignments / labs (50%): weighted according to effort**
- **Final project / lab (20%)**
- **In-class assignments (bonus 5%)**
- **Final (letter) grades will be based on a slight modification of a straight scale**
 - You pass if you score >50%

Topics

Week #	Topics
	Introduction
1	Bits and Bytes
	Integers
2	Machine Prog: Basics
	Machine Prog: Control
3	Machine Prog: Procedures
4	C language
	More C language
	The Memory Hierarchy
5	Cache Memories
	Linking
6	ECF: Exceptions
7	ECF: Processes & Signals
8	System Level I/O
9	Virtual Memory
10	Dynamic Memory Allocation
11	Network Programming
12	Concurrent Programming
13	The Future of Computing
14	F1 demo and short presentations

*Welcome
and Enjoy!*